

Maximum Mean Discrepancy in High-Dimensional Biology

Frederick Boehm^{1,*}

11 May, 2021

Abstract

Text of abstract

Contents

1	Specific Aims	1
2	Background and Significance	2
3	Approach	3
4	Methods	3
5	Results	3
6	Discussion	3
7	Conclusion	3
8	Acknowledgements	3
9	Colophon	4

¹ University of Massachusetts Medical School

* Correspondence: [Frederick Boehm <frederick.boehm@gmail.com>](mailto:Frederick.Boehm@gmail.com)

Keywords: keyword 1; keyword 2; keyword 3

Highlights: These are the highlights.

1 Specific Aims

The volume and dimensionality of molecular biology data has exploded over the last twenty years. Following the sequencing of the human genome, technologies such as DNA sequencing, RNA sequencing, and single-cell RNA sequencing have proliferated. The new experimental designs and high dimension of modern biological data require parallel development of new statistical methods and computational tools for their analysis. Maximum mean discrepancy (MMD) is a new statistical tool with novel mathematical properties [5, 4, 3, 1]. We propose to explore its mathematical and statistical properties, to apply MMD methods in the analysis of high-dimensional biological data, and to build user-friendly R software tools that implement MMD methods.

Aim 1: Apply MMD to high-dimensional biological data

Two-sample problems arise in many areas of biomedical research. These include differential gene expression, comparability of data from different laboratories, and classification of cancers. Borgwardt et al. [1] examined

performance characteristics of four multivariate two-sample tests in a collection of four biological settings. They found that MMD outperformed the other three tests. settings of

Aim 2: Create and maintain an open source R package that implements MMD methods

Current R software for MMD methods doesn't provide significance thresholds. Gretton's website hosts Matlab software that implements the methods in Gretton et al. [5], Gretton et al. [3] and Gretton et al. [4]. We'll implement these methods in R and C++ before assembling the code into a well documented, user-friendly R package [7, 2].

2 Background and Significance

MMD is a statistical method for determining whether two high-dimensional samples arise from distinct distributions. Throughout this document, we use the term "sample" to refer to measurements on a collection of subjects from a single group. In introductory statistics courses, one often studies Student's t-test [8] when discerning whether two samples come from the same distribution. Implicit to the t-test is the assumption that the two samples arise from normal distributions with a single, common variance parameter.

In high-dimensional settings, the analog of Student's t-test is Hotelling's T^2 test [6]. Like its univariate counterpart, Hotelling's test assumes a normal (multivariate joint normal) distribution of the random errors and a shared covariance matrix for the two groups. Departures from these assumptions may give rise to type I and type II errors in testing.

To combat limitations of existing high-dimensional, two-sample tests, Gretton et al. [4] proposed a kernel-based test for the two-sample problem. Drawing on reproducing kernel Hilbert space theory, they formulate their test as a measure of distance between two embeddings of probability measures in a RKHS. Gretton et al. [4] then write the test as a maximization over a subspace \mathcal{F} of their RKHS, $\mathcal{H}((1))$.

$$MMD[\mathcal{F}, p, q] = \sup_{f \in \mathcal{F}} (\mathbb{E}_x f(x) - \mathbb{E}_y f(y)) \quad (1)$$

Table 1: Microarray cross-platform comparability

platform	null_hypothesis	MMD	t_test	wolfowitz	smirnov
same	accepted	100	100	93	95
same	rejected	0	0	7	5
different	accepted	0	95	0	29
different	rejected	100	5	100	71

Table 2: Cancer diagnosis

health_status	null_hypothesis	MMD	t_test	wolfowitz	smirnov
same	accepted	100	100	97	98
same	rejected	0	0	3	2
different	accepted	0	100	0	38
different	rejected	100	0	100	62




Figure 1: A plot of random numbers

3 Approach

4 Methods

5 Results

Figure 1 shows how we can have a caption and cross-reference for a plot

Here is an example of inline code 3.14 in the middle of a sentence.

6 Discussion

7 Conclusion

8 Acknowledgements

9 Colophon

This report was generated on 2021-05-11 17:32:42 using the following computational environment and dependencies:

```
#> - Session info -----
#> setting value
#> version R version 4.0.4 (2021-02-15)
#> os      Ubuntu 21.04
#> system  x86_64, linux-gnu
#> ui      X11
#> language (EN)
#> collate en_US.UTF-8
#> ctype   en_US.UTF-8
#> tz      America/New_York
#> date    2021-05-11
#>
#> - Packages -----
#> package      * version      date      lib source
#> assertthat    0.2.1        2019-03-21 [1] CRAN (R 4.0.1)
#> backports     1.2.1        2020-12-09 [1] CRAN (R 4.0.3)
#> bookdown      0.21         2020-10-13 [1] CRAN (R 4.0.3)
#> cachem        1.0.4        2021-02-13 [1] CRAN (R 4.0.3)
#> callr         3.5.1        2020-10-13 [1] CRAN (R 4.0.3)
#> checkmate     2.0.0        2020-02-06 [1] CRAN (R 4.0.1)
#> cli           2.5.0        2021-04-26 [1] CRAN (R 4.0.4)
#> colorspace    2.0-1        2021-05-04 [1] CRAN (R 4.0.4)
#> commonmark    1.7          2018-12-01 [1] CRAN (R 4.0.1)
#> crayon        1.4.1        2021-02-08 [1] CRAN (R 4.0.3)
#> DBI           1.1.1        2021-01-15 [1] CRAN (R 4.0.3)
#> desc          1.2.0        2018-05-01 [1] CRAN (R 4.0.1)
#> devtools      2.3.2        2020-09-18 [1] CRAN (R 4.0.3)
#> digest        0.6.27       2020-10-24 [1] CRAN (R 4.0.3)
#> dplyr         1.0.6        2021-05-05 [1] CRAN (R 4.0.4)
#> ellipsis      0.3.2        2021-04-29 [1] CRAN (R 4.0.4)
#> evaluate      0.14         2019-05-28 [1] CRAN (R 4.0.1)
#> fansi         0.4.2        2021-01-15 [1] CRAN (R 4.0.3)
#> fastmap       1.1.0        2021-01-25 [1] CRAN (R 4.0.3)
#> fs            1.5.0        2020-07-31 [1] CRAN (R 4.0.3)
#> generics      0.1.0        2020-10-31 [1] CRAN (R 4.0.3)
#> ggplot2       3.3.3        2020-12-30 [1] CRAN (R 4.0.3)
#> glue          1.4.2        2020-08-27 [1] CRAN (R 4.0.3)
#> gt            0.1.0        2021-05-11 [1] Github (rstudio/gt@eff3be7)
#> gtable        0.3.0        2019-03-25 [1] CRAN (R 4.0.1)
#> highr         0.8          2019-03-20 [1] CRAN (R 4.0.1)
#> htmltools     0.5.1.1      2021-01-22 [1] CRAN (R 4.0.3)
#> knitr         1.31         2021-01-27 [1] CRAN (R 4.0.3)
#> lifecycle     1.0.0        2021-02-15 [1] CRAN (R 4.0.3)
#> magrittr      * 2.0.1       2020-11-17 [1] CRAN (R 4.0.3)
#> memoise       2.0.0        2021-01-26 [1] CRAN (R 4.0.3)
#> munsell       0.5.0        2018-06-12 [1] CRAN (R 4.0.1)
#> pillar        1.6.0        2021-04-13 [1] CRAN (R 4.0.4)
#> pkgbuild      1.2.0        2020-12-15 [1] CRAN (R 4.0.3)
#> pkgconfig     2.0.3        2019-09-22 [1] CRAN (R 4.0.1)
#> pkgload       1.2.0        2021-02-23 [1] CRAN (R 4.0.3)
```

```

#> prettyunits      1.1.1      2020-01-24 [1] CRAN (R 4.0.1)
#> processx         3.4.5      2020-11-30 [1] CRAN (R 4.0.3)
#> ps               1.5.0      2020-12-05 [1] CRAN (R 4.0.3)
#> purrr            0.3.4      2020-04-17 [1] CRAN (R 4.0.1)
#> R6               2.5.0      2020-10-28 [1] CRAN (R 4.0.3)
#> remotes          2.2.0      2020-07-21 [1] CRAN (R 4.0.3)
#> rlang            0.4.11.9000 2021-05-11 [1] Github (r-lib/rlang@7cd1f5c)
#> rmarkdown        2.7        2021-02-19 [1] CRAN (R 4.0.3)
#> rprojroot        2.0.2      2020-11-15 [1] CRAN (R 4.0.3)
#> sass             0.3.1.9003 2021-05-11 [1] Github (rstudio/sass@6166162)
#> scales           1.1.1      2020-05-11 [1] CRAN (R 4.0.1)
#> sessioninfo      1.1.1      2018-11-05 [1] CRAN (R 4.0.1)
#> stringi          1.6.1      2021-05-10 [1] CRAN (R 4.0.4)
#> stringr          1.4.0      2019-02-10 [1] CRAN (R 4.0.1)
#> testthat         3.0.2      2021-02-14 [1] CRAN (R 4.0.3)
#> tibble           3.1.1      2021-04-18 [1] CRAN (R 4.0.4)
#> tidyr            1.1.3      2021-03-03 [1] CRAN (R 4.0.3)
#> tidyselect       1.1.1      2021-04-30 [1] CRAN (R 4.0.4)
#> usethis          2.0.1.9000 2021-02-15 [1] Github (r-lib/usethis@aaf79d8)
#> utf8             1.2.1      2021-03-12 [1] CRAN (R 4.0.3)
#> vctrs            0.3.8.9000 2021-05-11 [1] Github (r-lib/vctrs@52157b7)
#> withr            2.4.2      2021-04-18 [1] CRAN (R 4.0.4)
#> xfun             0.22       2021-03-11 [1] CRAN (R 4.0.3)
#> yaml            2.2.1      2020-02-01 [1] CRAN (R 4.0.1)
#>
#> [1] /home/fred/R/x86_64-pc-linux-gnu-library/4.0
#> [2] /usr/local/lib/R/site-library
#> [3] /usr/lib/R/site-library
#> [4] /usr/lib/R/library

```

The current Git commit details are:

```

#> Local:   main /home/fred/work/professional-development/mmdproject
#> Remote:  main @ origin (https://github.com/fboehm/mmdproject.git)
#> Head:    [bda2491] 2021-05-10: added text to spec aims section

```

References

- [1] Karsten M Borgwardt et al. “Integrating structured biological data by kernel maximum mean discrepancy”. In: *Bioinformatics* 22.14 (2006), e49–e57.
- [2] Dirk Eddelbuettel and Romain François. “Rcpp: Seamless R and C++ Integration”. In: *Journal of Statistical Software* 40.8 (2011), pp. 1–18. DOI: [10.18637/jss.v040.i08](https://doi.org/10.18637/jss.v040.i08). URL: <https://www.jstatsoft.org/v40/i08/>.
- [3] Arthur Gretton et al. “A Fast, Consistent Kernel Two-Sample Test”. In: *NIPS*. Vol. 23. 2009, pp. 673–681.
- [4] Arthur Gretton et al. “A kernel method for the two-sample-problem”. In: *Advances in neural information processing systems* 19 (2006), pp. 513–520.
- [5] Arthur Gretton et al. “A kernel two-sample test”. In: *The Journal of Machine Learning Research* 13.1 (2012), pp. 723–773.
- [6] Harold Hotelling. “The Generalization of Student’s Ratio”. In: *The Annals of Mathematical Statistics* 2.3 (1931), pp. 360–378.
- [7] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. Vienna, Austria, 2021. URL: <https://www.R-project.org/>.
- [8] Student. “The probable error of a mean”. In: *Biometrika* (1908), pp. 1–25.